

1. **Understanding the [Research Repository](#) Structure** - By exploring the site, I identified the ability to use its search functionality and corresponding API to generate precise URLs for [topics of interest](#). This approach enabled efficient scraping by allowing me to target only relevant articles. To include abstracts and ensure a more comprehensive text dataset, I also decided to scrape individual article pages. I avoid duplicates by checking for unique article titles.
2. **Examining Site HTML and Identifying Tags** - Using Google Chrome's developer tools, I inspected the site's relatively static HTML structure and identified key elements containing metadata. I also accounted for pagination to ensure complete data collection.
3. **Scraping the Repository** - I employed Python's *requests* and *BeautifulSoup* libraries to first scrape the targeted search URLs, extracting the individual article URLs. For each article, I scraped metadata fields such as titles, abstracts, and other relevant content, storing the data in dictionaries. These were subsequently converted into a *pandas* DataFrame and exported as a structured CSV file.
4. **Keyword Frequency Analysis** - To analyze the text, I concatenated all article titles and abstracts into a single corpus and lowercased everything. Using Python's *re* package, I removed punctuation and digits, tokenized the text with *NLTK*, and filtered out stop words using *NLTK's* stopwords list. Finally, I performed word counts with *pandas* and visualized the results using the *plotnine* library, creating bar charts to showcase the most frequent keywords.